

Pacgate AI - Two-AIPC Deployment Handbook

Clone the repo on each machine, run the same install steps, and both machines become fully operational with deer-flow research and qm collaboration. Targets release 0.1.23 — handbook updated 2026-10-05 Chinese version: [AIPC-DEPLOYMENT-HANDBOOK-ZH.md](#) Prerequisites: Docker Desktop, Ollama, Node.js 24+. `install.ps1` pulls the models listed in `ollama-models.txt`.

⚠ Significant findings (2026-09-02) — read before deploying AIPC #2

These were discovered during the AIPC #1 pilot and are already fixed in this repo. AIPC #2 must pull the updated code (see Stage 1) so it gets these fixes.

Update 2026-09-23 (supersedes the 2026-09-15 "clone either" note). The two repos are NO LONGER identical — `pacgate-ai/pacgate-ai-pr` is 26 commits behind and missing the workflow-wiring fixes (`b7fc540`, `039afdc`), so a fork clone serves 10 built-in workflows instead of the firm's 222, with no error shown anywhere. Clone `JZKK720/pacgate-ai-pr` (see Stage 1). Which repo you push a release tag to still decides the GHCR namespace — see `plans/012-master-release-namespace.md`.

1. deer-flow agent could not query pacgate's legal databases. Root cause: no tool was wired to pacgate-api's `/api/kb/search` (RAG) or `/api/search` (legal connectors), and the memory adapter silently fell back to local files (`PacgateMemoryStorage` requires `PACGATE_MATTER_ID`). Fix: a new `pacgate-mcp` service (FastMCP) exposes `pacgate_kb_search`, `pacgate_connector_search`, `pacgate_list_connectors` to deer-flow. Registered in `deer-flow-extensions-config.json` beside `openviking`.
2. openviking MCP had the wrong API key baked in. `deer-flow-extensions-config.json` used the app key (`OPENVIKING_API_KEY`) but openviking's `root_api_key` is `OPENVIKING_ROOT_API_KEY`. A wrong key made openviking return 401, which rolled back the entire MCP tool load (deer-flow uses `asyncio.gather`), so no MCP tools appeared. Fix: use `OPENVIKING_ROOT_API_KEY` (template now uses `${OPENVIKING_ROOT_API_KEY}`).
3. `docker compose up -d --force-recreate deer-flow` wipes deer-flow's local DB. The SQLite DB, admin user, threads, and `.jwt_secret` live at `/app/backend/.deer-flow/` inside the container (not mounted). Recreating the container loses them → the frontend gets 401 and the `/setup` page appears. Use `docker compose restart deer-flow` for config changes; only recreate if you accept losing the local DB (then re-run `/setup`).
4. QM sign-in needs an email transport, not Outlook SMTP. The old SMTP path (`smtp.office365.com` + app password) is broken — Microsoft retired Basic Auth / app passwords for Exchange Online (Sep 2025). `qm check` failed with `535 5.7.139 Authentication unsuccessful`. Fix: qm's auth broker now supports two transports — Mailpit SMTP catcher (pilot; links land at `http://localhost:8025`, no key needed) and Resend (production; supply `RESEND_API_KEY`). See Stage 4. (Originally this item said Resend was the only option; Stage 4's two-transport setup supersedes it.)
5. The qm web-ui cannot self-authenticate. Its server (`/app/server/index.ts`) sets `AUTH_MODE = COOKIE_AUTH ? "dev" : "portal"`. Because `CORE_SIGNING_SECRET` is set, it's in portal mode and requires a portal-issued identity token. There is no no-secret way to reach the web-ui directly — you must run `portal+auth` (Resend) or an external OIDC provider. `ADMIN_GRANTS` is an authorization seed, not a sign-in.
6. Git push to `JZKK720/pacgate-ai-pr` was blocked for the `pacgate-ai` account (403, needs 2FA grant). Workaround: the `pacgate-ai` account could fork and push there. All fixes are now merged into both `main` branches, so this is historical — but note the two remotes are separate publishing targets for GHCR (see `plans/012-master-release-namespace.md`).

Architecture: two identical machines

Both AIPCs run the complete stack:

Each AIPC machine:

```
nginx :8089 -> pacgate-api :8080 (Rust metadata API)
          -> deer-flow :8001 (research workspace)
Postgres :5432 (local metadata DB)
OpenViking :1933 (long-term memory lane, MCP)
qm :8182 (co-working workspace, runs via `qm up`)
Ollama :11434 (native, GPU/NPU)
```

Each machine is self-contained and independently operational. Lawyers on either machine can use both research mode (deer-flow at <http://localhost:8089/research/>) and collaboration mode (qm at <http://localhost:8182>) without depending on the other machine.

If you later want shared matter data across both machines, connect them with a private mesh (Tailscale or WireGuard) and decide on a sync or single-authority model. That is a post-pilot decision, not a deployment prerequisite.

What you need before starting

- GitHub access to the source repo — **JZKK720/pacgate-ai-pr** (public; a plain clone needs no auth). Do not clone the **pacgate-ai** fork — it is 26 commits behind and serves 10 workflows instead of 222 (see Stage 1). A PAT or `gh auth login` is only required if you intend to push.
- Docker Desktop running on both AIPCs
- Ollama running on both AIPCs (`install.ps1` pulls the models it needs)
- `ollama signin` completed on each AIPC if the cloud-tagged deepseek models are in use
- Node.js 24+ installed on both AIPCs (for qm)
- No `docker login ghcr.io` needed — the Pacgate runtime images are published as public GHCR packages (see Stage 0).

Stage 0: Runtime images (dev machine, already done)

The runtime is published on GHCR and needs no rebuild on the AIPC.

Current versions (the ones to pull):

Image	Status
ghcr.io/jzkk720/pacgate-api:0.1.23	Published, public. Adds the matter-workspace rollup (GET /api/matters/:id/workspace).
ghcr.io/jzkk720/pacgate-mcp:0.1.23	Published, public. Exposes 19 MCP tools to deer-flow (adds <code>pacgate_get_workspace</code> , <code>pacgate_read_memory</code> , <code>pacgate_write_memory</code>).
ghcr.io/jzkk720/deer-flow-pacgate:0.1.23	Published, public.
ghcr.io/jzkk720/deer-flow-frontend-pacgate:0.1.23	Published, public.
ghcr.io/jzkk720/ocr-service:0.1.23	Published, public. PaddleOCR extraction; first-class since 0.1.16.
ghcr.io/volcengine/openviking@sha256:46f9e34c...	Pinned by digest in <code>compose.prod.yaml</code> . Upstream public image.

Namespace and version corrected 2026-09-22. This table previously listed `ghcr.io/pacgate-ai/*` at 0.1.0/0.1.3. `pacgate-ai` is the legacy mirror; the live namespace is `jzkk720` and publishing moved there in plan 016. The old `pacgate-ai/...-frontend-pacgate:0.1.0` row was also labelled "Published" while returning 404 - it never resolved. All five `jzkk720/*` images at the current pin (0.1.23, table above) return HTTP 200 anonymously.

Historical release table (retained for provenance, superseded):

Image	Status
ghcr.io/pacgate-ai/pacgate-api:0.1.3	Published at the time. Fixed the 0.1.1 container-networking bug.
ghcr.io/pacgate-ai/pacgate-mcp:0.1.3	Published at the time.
ghcr.io/pacgate-ai/deer-flow-pacgate:0.1.3	Published at the time.
ghcr.io/pacgate-ai/deer-flow-frontend-pacgate:0.1.0	Never published - 404. Do not use.

All Pacgate packages must be set to public visibility on GHCR so an AIPC can pull without registry credentials. Verify before rollout:

```
# Expect HTTP 200 with no docker login. 401/403 means the package is still private.
# 404 means the tag does not exist in that namespace – usually a drift between the
# pins in compose.prod.yaml and where the images were actually published.
#
# (The Accept header is REQUIRED – omit it and a public manifest returns 404, not 200.)
#
# This reads the four pins straight out of compose.prod.yaml, so it can never go
# stale: previously this snippet hard-coded an old tag that had been removed, and
# reported a healthy system as broken.
$acc = "application/vnd.oci.image.index.v1+json,application/vnd.docker.distribution.manifest.list.v2+json,ap
$compose = "deploy/client-bundle/compose.prod.yaml"
$pins = Select-String -Path $compose -Pattern "image:\s*(ghcr\.io/[^\s]+)" -AllMatches |
    ForEach-Object { $_.Matches } | ForEach-Object { $_.Groups[1].Value } |
    Where-Object { $_ -notmatch "openviking" } | Sort-Object -Unique

foreach ($pin in $pins) {
    $repo = $pin -replace "^\ghcr\.io/", ""          # owner/name:tag
    $name = ($repo -split ":")[0]                  # owner/name
    $tag = ($repo -split ":")[1]
    $t = (Invoke-RestMethod "https://ghcr.io/token?scope=repository:$name:pull").token
    $code = (Invoke-WebRequest -Uri "https://ghcr.io/v2/$name/manifests/$tag" `
        -Headers @{Authorization="Bearer $t"; Accept=$acc} `
        -Method Head -UseBasicParsing).StatusCode
    "{0,-58} => HTTP {1}" -f $pin, $code
}
```

To flip it (GitHub web UI — the API route 404s for personal accounts): GitHub → your profile → Packages → **pacgate-api** → Package settings → Visibility → Public → Save. Repeat for **deer-flow-pacgate**. This is safe: the images contain only the compiled binary and SQL migrations, every secret is injected at runtime via `.env`, and the installer already has full source access to the same code.

Only rebuild and push if the Rust source changes, from the dev machine. Take the tag from the compose pin rather than typing a version — a hard-coded tag here has gone stale twice, and one stale copy already reported a healthy system as broken:

```
cd c:\Users\cubeccloud-io\github-pr\pacgate-ai-pr
# Pattern reads the CURRENT namespace (jzkk720). The old form matched
# 'pacgate-ai/pacgate-api' and no longer matches any line in compose.prod.yaml,
# so $tag came back EMPTY and the build/push below silently used a blank tag.
$tag = (Select-String -Path deploy/client-bundle/compose.prod.yaml `
    -Pattern 'jzkk720/pacgate-api:(\S+)').Matches.Groups[1].Value
if (-not $tag) { throw 'could not read the image tag from compose.prod.yaml' }
docker build -t ghcr.io/jzkk720/pacgate-api:$tag -f pacgate-ai/Dockerfile ./pacgate-ai
docker push ghcr.io/jzkk720/pacgate-api:$tag
```

In practice prefer the `build-ghcr.yml` workflow so all four images stay in step — see `deploy/README-BUILD.md` and `plans/012-master-release-namespace.md`.

Do not rebuild on the AIPC — the pilot runs the published digests.

Port note: the stack binds `nginx` to host port `8089` (the value committed in `deploy/client-bundle/compose.prod.yaml`). If that port is already in use on the machine, edit the `ports:` entry for `nginx` and use the new port in all verification URLs below.

Stage 1: Clone the repo on each AIPC

On both machines:

```
cd C:\
git clone https://github.com/JZKK720/pacgate-ai-pr.git
cd pacgate-ai-pr
git remote -v # origin MUST be JZKK720/pacgate-ai-pr
```

AIPC #2 note (CORRECTED 2026-09-23): the two repos are NO LONGER IDENTICAL, so "either clone works" is no longer true. Clone `JZKK720`.

`pacgate-ai/pacgate-ai-pr` is 26 commits behind (last checked 2026-09-23) and is missing `b7fc540` and `039afdc`, so it still carries the original workflow-wiring defect. The 15 workflow YAMLS are present in the fork but are not wired into `pacgate-api`, so the API serves 10 built-in workflows instead of the firm's 222 — with no error shown anywhere. The identity claim above was accurate when written; the divergence came afterwards, which is exactly why a doc must not assert two things are in sync without a check that can fail.

The namespace difference is unchanged — which repo you push a release tag to decides which GHCR namespace the images publish into. See `plans/012-master-release-namespace.md`.

Cloning needs no credentials now that the repos are public; a PAT or `gh auth login` is only required to push.

If the repo is private and GitHub prompts for credentials, use a personal access token or the GitHub CLI (`gh auth login`).

Stage 2: Deploy the core stack (both machines, identical steps)

Run these steps on each AIPC. The Docker Compose stack starts `pacgate-api`, Postgres, `nginx`, and `deer-flow`.

```
cd C:\pacgate-ai-pr\deploy\client-bundle
copy .env.example .env
notepad .env
```

Fill in these values:

```
PACGATE_DB_PASSWORD=<generate a strong password>
PACGATE_JWT_SECRET=<generate a random hex string>
PACGATE_TENANT_ID=pacgate-law
OPENVIKING_ROOT_API_KEY=<generate a 32-char hex string>
OPENVIKING_API_KEY=<generate a 32-char hex string>
```

`OPENVIKING_API_KEY` is required — the installer renders `deer-flow-extensions-config.json` from it and stops with an error if it is missing or left as `change-me`.

Generate secrets if you need them:

```
# DB password (16 hex)
-join ((1..16) | ForEach-Object { '{0:x}' -f (Get-Random -Maximum 16) })

# JWT secret (32 hex)
-join ((1..32) | ForEach-Object { '{0:x}' -f (Get-Random -Maximum 16) })

# OpenViking keys (32 hex each) — generate a fresh value per line
```

```
-join ((1..32) | ForEach-Object { '{0:x}' -f (Get-Random -Maximum 16) })
```

Run the installer:

```
.\install.ps1
```

The installer pulls the (public, no-login) GHCR images, renders **OPENVIKING_CONF_CONTENT** and **deer-flow-extensions-config.json** from the **.env** secrets, starts the Docker Compose stack, and pulls the Ollama models listed in **ollama-models.txt**. If models are already pulled, this step is fast.

Verify the core stack:

```
docker compose -f compose.prod.yaml ps
curl http://localhost:8089/version
curl http://localhost:8089/pacgate/health
```

Expected: all eight containers running (pacgate-db, pacgate-api, deer-flow, deer-flow-frontend, pacgate-mcp, ocr-service, openviking, nginx); **/version** returns **{"version":"0.1.23","revision":"<git sha>"}** and **/pacgate/health** returns **ok**.

Do not probe **/health** at the nginx root. nginx routes **/** to the deer-flow frontend by design, so **curl http://localhost:8089/health** returns the frontend's 404 page - that reads as a failure but is correct routing. **/version** is also at the root (mapped to the API's **/build-info**); only **/pacgate/*** paths reach the API.

Stage 3: Seed the tenant and provision accounts (both machines)

install.ps1 now does this automatically (step 6a). On a current install you do not need to run anything on this page - it is kept for recovery, and the commands below are what the installer runs. Check the install output for **[OK] tenant 'default-firm' present** and **[OK] admin '...' registered**.

How accounts work since 0.1.22 (read before provisioning people)

- **POST /api/auth/register** is first-user-only: it creates exactly one account on a fresh deployment (the bootstrap admin) and refuses afterwards with 403.
- Every account after the first is created by the admin through **POST /api/auth/users** (Bearer = the admin's token).
- Research workspace sign-in = email + password. Collaboration (qm) sign-in = one-time emailed links allowlisted via **AUTH_ALLOWED_EMAILS**.

On each machine, seed the default tenant, then let the installer bootstrap the admin:

```
# Seed the tenant (idempotent).
#
# THE SLUG MUST MATCH PACGATE_TENANT_ID. It defaults to "default-firm", and the
# registration below looks the tenant up by that slug. An earlier version of this
# page used 'pacgate-law', which does not match, so registration failed with:
#
# {"error":{"code":"internal_error",
#   "message":"default tenant not found: matter not found: row not found"}}
#
# That reads like a database fault and is really a naming mismatch. If you set
# PACGATE_TENANT_ID to something else in .env, use that value here instead.
docker exec pacgate-db psql -U pacgate -d pacgate -c "INSERT INTO tenants (name, slug) SELECT 'Default Firm'
```

Provision an attorney user (run by the admin once the admin account exists):

```
$login = Invoke-RestMethod -Uri "http://localhost:8089/pacgate/api/auth/login" -Method Post `
  -Body '{"email":"admin@pacgate-law.com","password":"<admin-password>"}' `
  -ContentType "application/json"
$body = @{"email"<attorney-email>; password=<attorney-password>; role="attorney"} | ConvertTo-Json
```

```
Invoke-RestMethod -Uri "http://localhost:8089/pacgate/api/auth/users" -Method Post `
  -Headers @{Authorization="Bearer $($login.token)"} -Body $body -ContentType "application/json"
```

Register the qm bridge service account (the installer does this; manual for recovery). Same admin-provisioned route, same shape as above, with `email="qm-bridge@pacgate-law.com"`.

For the qm collaboration surface: set `AUTH_ALLOWED_EMAILS` (comma-separated) and `ADMIN_GRANTS=<email>:org_admin` in `deploy/client-bundle/qm-pacgate/.env`, then re-run `setup-qm.ps1`.

Stage 3.5: Verify OpenViking memory service (both machines)

OpenViking is the long-term memory lane: deer-flow and qm store conversational context there and recall it in later sessions. It starts as part of the compose stack.

```
curl http://localhost:1933/health
```

Expected: `{"status":"ok","healthy":true,...}`. The installer renders the OpenViking config (Ollama embedding + VLM) into `.env` as `OPENVIKING_CONF_CONTENT` and seeds the server's `ov.conf` on first boot.

Functional check (optional, uses the root key from `.env`):

```
$key = (Get-Content .env | Select-String '^OPENVIKING_ROOT_API_KEY=').Line.Split('=')[1]
$body = '{"jsonrpc":"2.0","id":1,"method":"tools/list"}'
curl.exe -s -X POST http://localhost:1933/mcp -H "X-API-Key: $key" -H "Content-Type: application/json" -H "A
```

Expected: a tool list including `find`, `search`, `read`, `remember`.

Boundary rule: OpenViking stores conversational context only (decisions, preferences, working knowledge). Matter documents and T1-T4-controlled content stay in `pacgate-api/pacgate-rag`.

Stage 4: Bootstrap qm (both machines, identical steps)

qm runs separately from the Docker Compose stack. Bootstrap it on each machine after the core stack is healthy.

```
cd C:\pacgate-ai-pr\deploy\client-bundle
.\setup-qm.ps1
```

The script prompts for: - Administrator work email (lowercased) - Pacgate bridge email:

`qm-bridge@pacgate.local` - Pacgate bridge password: the one you registered in Stage 3

The script generates signing secrets, creates `.env` in the `qm-pacgate` directory, validates the config with `qm check`, builds the sandbox image with `qm sandbox build`, fetches the static docker CLI the core shells (SHA-256 verified), and builds the `qm-pacgate-sandbox-local` exec-daemon wrapper image.

QM sign-in needs an email transport. The qm auth broker delivers sign-in one-time links. Two supported transports:

- Local/pilot topology (Mailpit SMTP catcher): `SMTP_HOST=mailpit SMTP_PORT=1025 AUTH_EMAIL_TRANSPORT=smtp SMTP_TLS=none`. Links LAND IN THE MAILPIT INBOX — open `http://localhost:8025` to pick up a sign-in link in a pilot. No real email is sent.
- Production topology (Resend): `AUTH_EMAIL_TRANSPORT=resend` with `RESEND_API_KEY` in `deploy/qm-pacgate/.env`, and `AUTH_EMAIL_FROM` set to a Resend-verified sender (an Outlook address is not verified; Microsoft retired Basic Auth/app passwords for Exchange Online, so raw Outlook SMTP is NOT a supported transport):

```
RESEND_API_KEY=re_XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
AUTH_EMAIL_FROM="PacGate <onboarding@resend.dev>"
```

Start qm:

```
cd C:\pacgate-ai-pr\deploy\client-bundle\qm-pacgate
```



```
docker compose -f compose.qm.yaml up -d
```

Use compose, NOT `qm up`. The `@yc-software/qm` CLI's docker lifecycle is POSIX-only - its `which()` shells `/bin/sh`, which does not exist on Windows (measured 2026-10-06: `execFileSync("/bin/sh", ...)` → `ENOENT`), so `qm up`, `qm down`, and `qm status` die with "docker not found on PATH" on every AIPC. `compose.qm.yaml` describes the same topology (same names, volumes, network) and additionally carries the Pacgate patches: `patch/pi-models.ts` and `patch/local-sandbox.ts` are bind-mounted over the core's source, the static docker CLI and the host docker socket are mounted, and the sandbox wrapper image is wired. Those mounts survive every `down/up -d` cycle - the model routing and the sandbox lane are durable through compose. Do NOT run `docker compose -f compose.qm.yaml up -d` and a `qm up` against the same directory (they would fight over the same named volumes and network).

Verify qm:

```
# Open http://localhost:8182 (web-ui) — requires a portal identity token
# Open http://localhost:8181 (portal) — the sign-in front door
# Sign in with the admin email (magic link via the configured transport; Mailpit in pilot)
# Send a test message
# Ask: "List available pacgate workflows"
```

Web-ui auth reality: the qm web-ui (:8182) cannot authenticate anyone on its own — it must be reached through the portal (:8181), which issues the identity token. If you open :8182 directly you'll see "reached through the portal." Always go through :8181.

Dev-mode alternative (pilot / single-user): for a local pilot you can run qm in dev/cookie mode instead of the portal. Recreate `qm-pacgate-core` with `NODE_ENV=development + ALLOW_UNAUTHENTICATED_CORE=1` and no `CORE_SIGNING_SECRET`, and `qm-pacgate-web-ui` with no `CORE_SIGNING_SECRET`. Then `POST /signin` works directly at :8182 with `{"user": "<principal>"}` and no Resend key is needed. This is not production-correct (no auth) — use it only for a single-user pilot.

Stage 5: Verify deer-flow (both machines)

On each machine, verify the research workspace:

```
# Open http://localhost:8089/research/
# Select or create a matter
# Ask: "Summarize recent force majeure case law in China"
# Verify: response includes citations
# Verify: response is saved to matter memory
```

Stage 5.5: Full-loop operations — deer-flow ↔ QM ↔ OpenViking

Both workspaces share one OpenViking long-term memory lane and one pacgate-api metadata store. This section documents how the two systems talk to each other and how to verify the loop end-to-end.

Topology (verified 2026-09-04)

```
pacgate-ai-bundle_default (Docker Compose network)
■■■■ openviking :1933 ← long-term memory (MCP: find/search/read/remember)
■■■■ pacgate-api :8080 ← metadata API (matters/workflows/connectors)
■■■■ pacgate-mcp :8000 ← FastMCP bridge exposing pacgate KB/connector search
■■■■ deer-flow :8001 ← research workspace (consumes openviking + pacgate-mcp)
■■■■ nginx :8089 ← ingress (deer-flow frontend + /pacgate/ API)

qm-pacgate (qm up network)
■■■■ qm-pacgate-core ← co-working agent runtime
■■■■ qm-pacgate-web-ui ← browser chat UI (:8182)
■■■■ qm-pacgate-pg ← qm Postgres
```

Bridge: `qm-pacgate-core` is ALSO joined to `pacgate-ai-bundle_default`, so it can

reach `openviking:1933`, `pacgate-api:8080`, and `host.docker.internal:11434` (Ollama).

How the loop works

1. `deer-flow` → OpenViking: `deer-flow-extensions-config.json` registers `openviking` as an HTTP MCP server at `http://openviking:1933/mcp` with the root API key (`OPENVIKING_ROOT_API_KEY`). Research runs store and recall conversational context there.
2. `deer-flow` → `pacgate`: `pacgate-mcp` (FastMCP) exposes `pacgate_kb_search` / `pacgate_connector_search` / `pacgate_list_connectors` to `deer-flow`, backed by `pacgate-api`'s RAG + legal connectors.
3. QM → OpenViking: the QM core has `OPENVIKING_URL=http://openviking:1933` and `OPENVIKING_API_KEY` (the root key). The `pacgate-qm` sandbox tool (`ov-remember` / `ov-search` / `ov-read`) calls OpenViking from inside the agent sandbox at `http://host.docker.internal:1933` with `OPENVIKING_ACCOUNT=pacgate-law` + `OPENVIKING_USER`.
4. QM → `pacgate`: the `pacgate-qm` sandbox tool logs into `pacgate-api` (`PACGATE_API_EMAIL` / `PACGATE_API_PASSWORD`) to discover workflows, bind a QM scope to a matter, read/write matter memory, and execute workflows.

Verify the full loop

```
# 1. OpenViking MCP responds (root key from .env)
$key = (Get-Content .env | Select-String '^OPENVIKING_ROOT_API_KEY=').Line.Split('=')[1]
$body = '{"jsonrpc":"2.0","id":1,"method":"tools/list"}'
curl.exe -s -X POST http://localhost:1933/mcp -H "X-API-Key: $key" -H "Content-Type: application/json" -H "A
# Expected: tools find/search/read/remember

# 2. QM core reaches openviking + pacgate-api + Ollama
docker exec qm-pacgate-core sh -c "curl -s -o /dev/null -w 'openviking:%{http_code}\n' http://openviking:1933

# 3. QM sign-in works (dev/cookie mode)
$body = '{"user":"admin@pacgate-law.com"}'
curl.exe -s -X POST http://localhost:8182/signin -H "Content-Type: application/json" -d $body
# Expected: {"ok":true,"user":"admin@pacgate-law.com"}
```

QM ↔ bundle network join (idempotent)

QM core is joined to the bundle network so it can resolve `openviking` and `pacgate-api` by name. If a `qm up/qm down` cycle drops the join, re-apply it:

```
docker network connect pacgate-ai-bundle_default qm-pacgate-core
```

Note: QM is intentionally not a Docker Compose service in the bundle. It is managed by the `@yc-software/qm` CLI (`qm up` / `qm down`) with its own lifecycle. The network join is the only coupling — keep it that way. Do not rewrite QM as compose services; that fights the `qm` CLI and would recreate the containers on every `qm up`.

Stage 6: Smoke test checklist (both machines)

Run this checklist on each AIPC independently.

Core stack

- `[] docker compose -f compose.prod.yaml ps` shows 5 services up (incl. `openviking`)
- `[] curl http://localhost:8089/version` returns the release version + revision JSON
- `[] curl http://localhost:8089/pacgate/health` returns `ok` (NOT `/health` at the root - see Stage 2)

- [] `curl http://localhost:1933/health` returns healthy JSON
- [] Postgres has the **pacgate-law** tenant
- [] Admin user can log in at `http://localhost:8089/api/auth/login`
- [] deer-flow returns a real research response at `http://localhost:8089/research/`

qm collaboration

- [] `npm exec qm -- status` shows qm running
- [] `http://localhost:8182` loads the qm web UI
- [] Admin can sign in
- [] qm can list Pacgate workflow categories
- [] qm can execute one Pacgate workflow through the bridge

Ollama

- [] `ollama list` shows the required models
- [] deer-flow can call Ollama for inference
- [] qm can call Ollama for inference

Data

- [] `./data/tenants/` directory exists and is writable
- [] `./openviking/` directory exists and persists across restarts
- [] Document upload works through the API
- [] Matter memory persists after a deer-flow research run
- [] Cross-session recall: a fact stored via OpenViking **remember** is recalled via **search** in a later session

Managing the stack after deployment

Start and stop

```
# Start core stack
docker compose -f compose.prod.yaml up -d

# Stop core stack
docker compose -f compose.prod.yaml down

# Start qm
cd C:\pacgate-ai-pr\deploy\client-bundle\qm-pacgate
docker compose -f compose.qm.yaml up -d

# Stop qm (NO -v: .env and the volumes stay)
docker compose -f compose.qm.yaml down
```

`qm up / qm down` cannot run on Windows (the CLI's `which()` shells `/bin/sh` → `ENOENT`); `compose` is the only working qm lifecycle on the AIPC and it carries the patches that keep model routing and the sandbox lane durable.

Update to a new version

```
cd C:\pacgate-ai-pr\deploy\client-bundle
.\install.ps1 -Update
```

-Update now does all of the following, so a separate **git pull** is no longer needed (it was the most easily forgotten step, and forgetting it silently ran new images against old config):

Step	What it does
1	Refreshes the repo working tree (fast-forward only)
2	Re-renders the MCP config from the template, backing up any change
3	Pulls new GHCR images
4	Restarts the stack, and reloads nginx

It refuses, rather than guessing, when the machine has local work:

- uncommitted changes in the repo -> the repo update is skipped, and the changed files are listed. Nothing is stashed, reset, or discarded.
- local commits the remote does not have -> it will not merge or rebase. The repo stays as-is and the rest of the update continues.

Use **-SkipRepoPull** if you deliberately want an image-only update:

```
.\install.ps1 -Update -SkipRepoPull
```

qm is reported, not auto-updated

If qm is running on the machine, **-Update** also checks whether its sandbox image still matches its source:

```
[OK] qm sandbox matches its source (38e062ec7c5c...)
```

or

```
[WARN] qm sandbox source has CHANGED since the image was pinned.
qm is running OLD skills and tools. Rebuild + repin:
  cd deploy\qm-pacgate
  npm exec qm -- sandbox build # then repin the printed digest
  pwsh -File ..\..\scripts\qm-sandbox-fingerprint.ps1 -Write
```

This matters because qm's agent executes inside a sandbox image pinned by digest in **deploy/qm-pacgate/qm.config.jsonc**. Digest pinning is correct - the isolation boundary should be immutable - but it means a repo update can change **deploy/qm-pacgate/sandbox/** while the image stays exactly as it was. The agent then keeps running the old skills and tools, with no error anywhere.

-Update reports this rather than rebuilding on its own, because the rebuild needs Node 24 + npm + buildx and the digest must be repinned afterwards - a config change we should not make unattended on a client machine. A wrong automatic rebuild is a worse failure than a visible warning.

Data is preserved across an update: - **./data/tenants/** (volume mount) - matters, documents, memory - Postgres data (named volume) - metadata database

Is this machine current? Ask it.

```
curl.exe -s http://localhost:8089/version
```

```
{"version":"<release>","revision":"<git sha>"}
```

This reports the version compiled into the running pacgate-api binary, and the commit it was built from. It deliberately does not echo the compose pin or the image tag: those record what was deployed, and the failure worth catching is exactly the case where the deployed artifact and the running process disagree.

Before this existed there was no way to tell a current machine from one behind - both looked identical from the outside, and the only check was to SSH in and read the compose file, hoping the containers matched it.

Switch models

deer-flow (research workspace): 1. Edit `deer-flow-config.yaml` - reorder the `models` list (first entry = default)
2. Restart: `docker compose -f compose.prod.yaml restart deer-flow`

qm (co-working workspace): 1. Edit `deploy/client-bundle/qm-pacgate/qm.config.jsonc` if the model set changes 2. Recreate core (compose keeps the patch mounts, so the routing stays durable): `powershell cd C:\pacgate-ai-pr\deploy\client-bundle\qm-pacgate docker compose -f compose.qm.yaml up -d --force-recreate core`

Register new users

`POST /api/auth/register` is first-user-only and returns 403 after the bootstrap account (see Stage 3). Create every later account through the admin route:

```
# /pacgate prefix required - see the note in Stage 3. Bearer = admin's token
# from POST /pacgate/api/auth/login (see Stage 3 for the login step).
$body = @{email="<user>@pacgate-law.com"; password="<password>"} | ConvertTo-Json
Invoke-RestMethod -Uri "http://localhost:8089/pacgate/api/auth/users" -Method POST `
  -Headers @{Authorization="Bearer <admin-token>"} -Body $body -ContentType "application/json"
```

Backup the database

```
docker exec pacgate-db pg_dump -U pacgate pacgate > backup.sql
```

Check logs

```
docker compose -f compose.prod.yaml logs -f pacgate-api
docker compose -f compose.prod.yaml logs -f deer-flow
```

Known limitations

- `qm up / qm down` do not work on Windows. The qm CLI's docker lifecycle is POSIX-only: its `which()` runs `execFileSync("/bin/sh", ...)`, which is `ENOENT` on every Windows AIPC (measured 2026-10-06). Use the compose path above (`docker compose -f compose.qm.yaml up -d`) - same topology, same container names, and it carries the Pacgate patches, so nothing is lost across recreates.
- Each machine has its own independent Postgres and `./data/tenants/` directory. Matter data is not shared between machines unless you later add a private mesh and a sync or single-authority model.
- The PkuLaw connector token is expired. Regenerate it at <https://mcp.pkulaw.com> and set `PKULAW_API_KEY` in `.env` if China-law search is needed during the pilot.
- Four WASM crates (citation-check, clause-parser, doc-validator, rule-engine) remain stubs. These are future-blueprint work and do not affect Phase 1 pilot functionality.
- Model selection: the API defaults to models that may not exist on the target machine. After Stage 3, apply per-tenant model overrides so the LLM tiers point at models actually present in `ollama list` on that machine. Recommended pilot set (benchmarked 2026-08-28): `gemma4:12b-it-qat` (Main — 13s/tool-round, schema-valid tool calls, verified end-to-end), `qwen3.8:27b-mtp-q4_K_M` (Mid — 73s/tool-round, stronger quality for batch tabular review), `nomic-embed-text:latest` (embeddings). Avoid reasoning-mode models (e.g. nemotron) for interactive tiers — they can hang long docx generations. See [plans/007-aipc-full-installation-handoff.md](#) Appendix A for the SQL template.

Files referenced

File	Purpose
<code>deploy/client-bundle/compose.prod.yaml</code>	Docker Compose for pacgate-api + deer-flow + Postgres + nginx

File	Purpose
deploy/client-bundle/install.ps1	One-click Windows installer for the core stack
deploy/client-bundle/setup-qm.ps1	qm bootstrap script (secrets, config, sandbox build)
deploy/client-bundle/.env.example	Template for client secrets
deploy/client-bundle/ollama-models.txt	Models to pre-pull
deploy/client-bundle/deer-flow-config.yaml	Multi-model deer-flow config (5 models, switchable)
deploy/qm-pacgate/qm.config.jsonc	qm local deployment config
deploy/SETUP-AND-OPERATIONS.md	Full 3-day on-site install guide (reference)
deploy/DEPLOYMENT-GUIDE.md	Engineer-level deployment details (reference)