

Pacgate-ai User Manual

For attorneys, paralegals, and partners at the law firm Version 0.3.0 — 2026-10-05 (matches Pacgate-ai release v0.1.23) 中文版：[USER-MANUAL-ZH.md](#) | PDF: [EN](#) / [ZH](#)


What Pacgate-ai is

Pacgate-ai is your firm's private legal AI assistant. It runs on the AI PC in your office — your documents, your matters, and your research history stay on that machine. There is no firm account on any cloud storage service. When you ask the assistant a question, the text of the conversation is processed by the AI model service your firm has enabled (for speed and quality); your document files themselves are never uploaded. See Privacy and security below for the exact split.

Pacgate-ai has two working modes, each for a different kind of task:

Mode	What it's for	When to use it
Research	Deep legal research with citations, document analysis, contract review	"Research this jurisdiction's recent case law on X"
Collaborate	Share documents, approve workflows, distribute work to the team	"Share this contract draft with the M&A team for review"

You switch between them by clicking a different link in your browser. Both run on the same AI PC; both see the same matters and documents.



How the pieces connect

Getting started

How to access

1. Open your web browser (Chrome, Edge, or Firefox)
2. Go to: <http://:8089> - Ask your IT person for the AI PC's IP address - If you're on the AI PC itself, use <http://localhost:8089>
3. You'll see the Pacgate-ai landing page with two options: - "Research a matter" → opens the research workspace - "Collaborate on a matter" → opens the collaboration workspace

Logging in

The two workspaces sign you in differently. Both give you a real, personal identity - your actions are attributed to you in the audit trail.

Research workspace (deer-flow) - username and password:

1. Click "Research a matter" on the landing page.
2. If you are not signed in yet, the app shows the login page.
3. Enter the username (usually your email) and the password your IT administrator gave you.
4. Your session stays active until you log out or close the browser.

There is no self-registration on the research workspace: accounts are created by your administrator, one person at a time. If you have no account yet, ask your IT administrator to provision one (they need only your name and your firm email).

Collaboration workspace (qm) - email sign-in link, no password:

1. Click "Collaborate on a matter" on the landing page.
2. Enter your firm email address on the sign-in page.
3. The system emails you a one-time sign-in link (valid for a short while; one use only).
4. Open the link from your inbox - it completes the sign-in and opens the workspace. There is no password to remember or reset.

Sign-in links only reach people on your administrator's allowlist. If "this address can't sign in" appears, you are not on the list yet - ask your administrator to add you.

In both workspaces, your identity shows at the edge of the screen (your email in the sidebar / account chip). Sign out from there when sharing a machine.

Mode 1: Research (deer-flow)

What this mode does

The research workspace is for tasks that need deep analysis:

- Legal research: "Find recent case law on force majeure clauses in China mainland courts, with citations"
- Contract review: "Analyze this contract for unfavorable clauses and risk areas"
- Document comparison: "Compare these two contract versions and show what changed"
- Tabular extraction: "Extract all payment terms from these 10 contracts into a table"
- Report generation: "Research cross-border M&A precedent in Hong Kong and produce a cited memo"

How to use it

1. Click "Research a matter" on the landing page
2. Select which matter you're working on (a dropdown of your firm's active matters)
3. Type your research question or task in the chat box
4. Press Enter or click Send
5. The AI will: - Break your request into steps - Search your firm's knowledge base and documents - Read and analyze the relevant files - Produce a structured response with citations to specific documents and pages
6. The response appears in the chat, with a panel on the right showing artifacts (files the AI produced)

Understanding the response



How an answer is generated

- Inline citations look like [1], [2] — click them to jump to the source document and page
- Artifacts (files produced by the AI) appear in the right panel — click to preview, download, or open
- Tool calls show what the AI did step-by-step (read document, search knowledge base, generate DOCX)

Working with documents

Upload a source document: - Drag a file into the chat, or click the attachment button - Supported: **.docx**, **.pdf**, **.txt**, **.md** - The document is stored under your matter's folder

Generate a new document: - Ask: "Generate a contract review memo for matter M-001 based on the research above" - The AI produces a **.docx** file via the document engine (Office documents via OfficeCLI; Word/Excel/PowerPoint all supported) - It appears as an artifact — click to download or preview

Convert between formats: - "Convert this document to Markdown" — the AI converts **.docx** / **.pdf** / **.pptx** / **.xlsx** / **.html** to clean Markdown text you can quote or edit - The reverse also works: ask for Markdown converted to **.docx** or **.pdf** for delivery - Conversion is local; the document content never leaves the AI PC during conversion

Read text out of a scanned document (OCR extraction): - Upload a scanned PDF or image of a document, then ask: "Read the text of this document" - The OCR engine reads the pages, records every text span with its page and position, and reports how complete the read was - If any page could not be parsed, the assistant says so plainly ("extraction incomplete") — a partial read is never presented as complete - Extraction results are kept per document version, so the same document is not re-read twice - Batch reading: "OCR every document in this matter" — the assistant works through a capped number of pages per run and reports the remaining budget

Sanitizing a document (preparing client material for outside processing):

This system can prepare a document so its client-identifying details are replaced before any outside model service processes it:

1. Ask: "Sanitize this document" — the sanitizer agent runs the redaction pipeline
2. The system replaces identifiers (ID numbers, phone numbers, names) with placeholders, records what was replaced in a sealed mapping, and reports the outcome plainly: verdict, how many replacements were made
3. The sanitized version is what may then be quoted in chat or sent for outside processing; the original stays behind
4. The mapping itself is not shown in chat; only an administrator can restore the original from it, and every restore is audit-logged
5. A document that fails the check is marked blocked and cannot leave the machine

Where your documents live (the matter workspace):

Every document — uploaded, generated, or shown to you by the AI — is filed under its matter:

Your matter

- Documents (each with versions and a sanitization status)
- Extraction records (what the OCR has read)
- Deliverables shown in chat (kept automatically, named like the file you saw)
- Search index (which documents the assistant can quote from)

To see it: ask the assistant "Show me this matter's workspace." To retrieve any document: ask for it by name ("download the memo we drafted"), preview it in the artifacts panel, or browse the matter document list in the UI. Downloads of

documents that carry client identity are held until their safety review passes — seeing "pending" on a fresh upload before its review is normal.

Version history: - Every document has versions (v1, v2, v3...) - Each edit or generation creates a new version - You can view and download any previous version

Research skills available

The research workspace comes with specialized legal skills:

Skill	What it does
Deep Research	Multi-step research with source gathering and cross-checking
Systematic Literature Review	Structured review of legal literature and precedent
Academic Paper Review	Analyze and summarize legal academic papers
Consulting Analysis	Business/legal consulting frameworks
GitHub Deep Research	Analyze a codebase or technical repository

The AI automatically selects the right skill based on your request. You don't need to choose manually.

Mode 2: Collaborate (qm)

What this mode does

The collaboration workspace is for working with your team:

- Share documents: "Share this contract draft with the M&A team"
- Approve workflows: Review and approve a document generation workflow before it runs
- Distribute work: "Send this research memo to all partners on matter M-001"
- Track approvals: See what's pending your approval and what you've approved

How to use it

1. Click "Collaborate on a matter" on the landing page
2. You'll see your firm's collaboration dashboard
3. The left sidebar shows your sessions — each is a conversation about a specific matter or topic
4. The center is the chat area — type messages or attach files
5. The right panel shows files shared in this session

Sharing documents

1. In a collaboration session, click the file/attachment button
2. Select a document from your matter (or upload a new one)
3. Choose who to share with: - Specific people (select attorneys from your firm) - The matter team (everyone on this matter) - The practice group (e.g., "all M&A attorneys")
4. Choose permission: View only or Can edit
5. Click Share

The people you shared with will see the document in their next collaboration session for that matter.

Approving workflows

When the AI proposes an action that needs approval (e.g., "I'll generate a contract review report and share it with the team"):

1. You'll see an approval card in the chat
2. Review what the AI wants to do
3. Click Approve to let it proceed, or Reject with feedback
4. The AI only acts after you approve

Ethical walls

Your firm's administrator can set ethical walls on specific matters. If a matter has an ethical wall:

- Only assigned attorneys can see that matter's documents
- The AI will not reference documents from a walled matter when working on other matters
- This prevents conflict-of-interest exposure

If you try to access a matter you're not assigned to, you'll see "Matter not found" — this is intentional, not an error.

Matters and documents

What is a matter?

A matter is a legal case, project, or engagement. Every document, research task, and collaboration session belongs to a matter. Matters belong to your firm (the tenant).

Your firm's matters might include: - M-001: Smith v. Jones litigation - M-002: Acme Corp M&A transaction - M-003: Patent application #2026-XXX

Document structure

Each matter has:

M-001 (matter)
■■■ Documents (versioned .docx, .pdf, .txt)
■■■ Uploads (source files you attach)
■■■ Memory (facts the AI extracted from your work)
■■■ Run history (what the AI did, when)

Document versions

Every document is versioned: - When the AI generates a document, it creates v1 - When you or the AI edits it, it creates v2, v3, etc. - You can download or preview any version - The current version is always the latest

The matter workspace (one view of everything)

Everything your matter holds is gathered in one place:

- Documents - every uploaded or generated document, with its sanitization state (pending / sanitized / blocked) and current version
- Extraction records - what the AI has read out of each document (page counts, whether any page failed), so you always know how complete the machine's view is
- Knowledge-base status - which documents are indexed for search and whether they have been through the safety pipeline

Ask the assistant in research mode: "Show me the workspace for this matter." It reads the aggregated view and reports what the matter holds before you start work - useful when picking up a matter someone else started.

Deliverables you are shown are saved for you

When the AI produces a document and shows it in the artifacts panel, the system automatically keeps a copy filed under the matter - named like the file you saw (for example **presented-report.docx**). That means:

- The deliverable you were shown remains available from the matter, not only in the chat that produced it
- Other people working on the matter can find it; it is not locked to your session
- It survives updates and restarts

If you want the same file with your own name on it, ask the assistant to upload it as a named document instead - both paths are available.

Citations

When the AI cites a source in its response: - **[1]** means it's referencing a specific document and page - Click the citation to open that document at the right page - The verbatim quote (≤ 25 words) is shown alongside the citation

This is critical for legal work — every AI claim should be traceable to a source document.

Personas

Pacgate-ai comes with legal personas — specialized AI assistants tuned for different practice areas:

Persona	Practice area
M&A Partner	Mergers and acquisitions
Litigation Associate	Court proceedings, discovery
IP Counsel	Intellectual property
Compliance Officer	Regulatory compliance
Cross-border Advisor	International/cross-border matters
...	20 total, firm-customizable

Your firm can customize personas or add new ones. Ask your administrator if you need a persona that doesn't exist yet.

The three-tier model system

Pacgate-ai uses three model tiers for different tasks:

Tier	Used for	Speed
Main	Deep analysis, contract review, complex reasoning	Slower, most capable
Mid	Tabular review, batch extraction, structured output	Medium
Low	Titles, short summaries, quick labels	Fastest

You don't choose the tier — the AI automatically picks the right model for each subtask. A quick label uses Low; a contract review uses Main.

Your firm's administrator configures which models serve each tier, and the set can change over time as better models become available. The research and collaboration chat assistants may use hosted models for speed; heavy document processing runs on the office AI PC.

Privacy and security



What stays in the building, and what is sent

What stays on your machine

- All documents — uploaded, generated, and edited

- All matter data — case files, research history, memory
- The knowledge base and search index
- All collaboration — shared within your firm's network only

What is sent to the AI model service

- The text of your conversation — your question and the specific passages the assistant reads to answer it are processed by the model service your firm has enabled, so answers come back fast and at high quality.
- Your document files themselves are never uploaded to that service, and no matter data is stored in any cloud account.

If a client matter requires zero external processing of any kind, tell your firm administrator before using the system on that matter.

Who can see what

- You can only see matters you're assigned to
- Documents are scoped to their matter
- Ethical walls prevent cross-matter data leakage
- Every document access and AI action is logged (audit trail)

Frequently asked questions

"The AI is slow"

Deep research tasks (Main tier) take longer because they run multiple steps: search, read, analyze, cross-check, generate. A full research report may take 2-5 minutes. Quick tasks (summaries, labels) take seconds. This is normal — quality takes time.

"I don't see a matter I should have access to"

Contact your firm administrator. They control matter assignments and ethical walls. If you're not assigned to a matter, it won't appear in your list.

"The AI cited something wrong"

Click the citation to open the source document. If the citation is inaccurate, tell the AI in the chat: "That citation is wrong — the correct page is 15, not 12." The AI will correct itself and update its memory for future tasks.

"Can I edit a document the AI generated?"

Yes. Download the **.docx** artifact, edit it in Word, and re-upload it. It becomes a new version (v2, v3...). The AI can also edit documents directly if you ask it to.

"How do I share my research with a colleague?"

1. In the research workspace, the AI produces an artifact (e.g., a research memo **.docx**)
2. Switch to the collaboration workspace
3. Start a session on the same matter
4. Share the document with your colleague

"What if the AI PC goes down?"

Your data is safe — it's on the AI PC's disk, in a Docker volume. When the AI PC restarts, Docker Desktop auto-starts the containers (if configured), and all your matters, documents, and history are preserved. If you need help, contact your IT administrator.

"How do I get a new version of the software?"

Your administrator will tell you when an update is available. They run a one-command update script. Your data is never affected by updates — only the software engine changes. All your matters, documents, and customizations are preserved.

Keyboard shortcuts

Shortcut	Action
Enter	Send message
Shift+Enter	New line in message
Ctrl+K	Search matters/documents (if available)
Esc	Close artifact preview

Getting help

- Technical issues (can't log in, page won't load): Contact your firm's IT administrator
 - Feature requests (want a new persona, new workflow template): Contact your firm's Pacgate-ai administrator
 - Training (how to use research or collaboration effectively): Ask your administrator to arrange a session with Cubecloud support
-

Glossary

Term	Meaning
Matter	A legal case, project, or engagement
Tenant	Your law firm (the organization boundary)
Artifact	A file produced by the AI (e.g., a .docx research memo)
Citation	A reference to a specific document and page, shown as [1], [2]
Persona	A specialized AI assistant for a practice area (e.g., "M&A Partner")
Ethical wall	A restriction preventing certain attorneys from seeing a matter's data
Tier	The model capability level (Main/Mid/Low) used for a task
Scope	The access boundary for a matter or practice group
Ollama	The local AI model runner on your AI PC